

# PUNJAB UNIVERSITY COLLEGE OF INFORMATION AND TECHNOLOGY



## **PROGRAMMING ASSIGNMENT # 1**

Course: Operating Systems

Submitted to: Sir Ahmad

Submitted by: Mariam Qadeem

Roll no: Bitf22m006

# QUESTION#1

In operating systems like Linux, the struct `task_struct` represents a *Process Control Block (PCB)*, which contains various details about a process, such as its state, ID, memory usage, and scheduling information. Each member in `task_struct` is specifically tailored to store an aspect of the process, making it easier to access and manage process information in a structured manner.

## Core Members of `task_struct`

### 1. `pid`

- **Code Representation:** `pid_t pid`
- **Description:** The unique identifier assigned to each process.

### 2. `state`

- **Code Representation:** `long state`
- **Description:** Represents the process's current state (e.g., running, sleeping, stopped, zombie).

### 3. `flags`

- **Code Representation:** `unsigned int flags`
- **Description:** Holds flags that provide metadata on the process, including its privileges or special conditions.

### 4. `parent`

- **Code Representation:** `struct task_struct *parent`
- **Description:** A pointer to the `task_struct` of the parent process, enabling hierarchy and tracking.

### 5. `children`

- **Code Representation:** `struct list_head children`
- **Description:** Manages a linked list of child processes, facilitating process tree structure.

## 6. sibling

- **Code Representation:** `struct list_head sibling`
- **Description:** A linked list node for navigating between sibling processes, representing other children of the same parent.

## 7. comm

- **Code Representation:** `char comm[TASK_COMM_LEN]`
- **Description:** Contains the name of the process's executable, used in various monitoring tools.

## 8. mm

- **Code Representation:** `struct mm_struct *mm`
- **Description:** Points to the memory management structure (`mm_struct`), managing the process's memory allocations.

### mm\_struct Members (Referenced by mm)

#### ➤ pgd

- **Code Representation:** `pgd_t *pgd`
- **Description:** Points to the Page Global Directory, which is essential for managing page tables and virtual memory.

#### ➤ mmap

- **Code Representation:** `struct vm_area_struct *mmap`
- **Description:** Pointer to a list of memory regions (`vm_area_struct`) allocated to the process.

#### ➤ total\_vm

- **Code Representation:** `unsigned long total_vm`
- **Description:** Tracks the total number of virtual memory pages the process has allocated.

#### ➤ rss

- **Code Representation:** `unsigned long rss`
- **Description:** The Resident Set Size (RSS), indicating how many physical memory pages are currently held by the process.

#### ➤ stack\_vm

- **Code Representation:** `unsigned long stack_vm`
- **Description:** Amount of virtual memory allocated specifically for the process's stack.

#### ➤ data\_vm

- **Code Representation:** `unsigned long data_vm`

- **Description:** The virtual memory size of the data section, which includes initialized and uninitialized data segments.

## 9. cred

- **Code Representation:** `const struct cred *cred`
- **Description:** Points to the credential structure, managing the user and group IDs, and security context.

## 10. files

- **Code Representation:** `struct files_struct *files`
- **Description:** Points to a structure that manages the open file descriptors for the process.

### files struct Members (Referenced by files)

#### ➤ fdt

- **Code Representation:** `struct fdtable *fdt`
- **Description:** The file descriptor table, managing open files for the process.

#### ➤ next\_fd

- **Code Representation:** `int next_fd`
- **Description:** The next available file descriptor in the process's file descriptor table.

## 11. fs

- **Code Representation:** `struct fs_struct *fs`
- **Description:** Points to filesystem information, including the current and root directories.

## 12. signal

- **Code Representation:** `struct signal_struct *signal`
- **Description:** Manages signal handling for the process, storing pending signals and signal masks.

### signal struct Members (Referenced by signal)

#### ➤ shared\_pending

- **Code Representation:** `struct sigpending shared_pending`

- **Description:** Stores signals that are pending and shared across threads.

➤ **rlim**

- **Code Representation:** `struct rlimit rlim[RLIM_NLIMITS]`
- **Description:** Array of resource limits, covering aspects like CPU time and memory usage limits.

### 13. cpu

- **Code Representation:** `int cpu`
- **Description:** Indicates the CPU on which the process is currently scheduled, aiding in load balancing.

### 14. start\_time

- **Code Representation:** `struct timespec start_time`
- **Description:** Records the start time of the process, useful in calculating the total running time.

### 15. exit\_code

- **Code Representation:** `int exit_code`
- **Description:** Holds the exit code provided by the process upon termination.

### 16. stack

- **Code Representation:** `void *stack`
  - **Description:** Points to the process's kernel stack, utilized during execution for storing function calls and local variables.
-

## QUESTION#2

### List of Linux system calls used in process creation and management

1. **fork:**  
Creates a new process by duplicating the calling process. The child process inherits the parent's memory space and gets a unique process ID (PID).
2. **vfork:**  
Similar to `fork`, but it creates a new process without copying the parent's address space until an `exec` call is made, improving performance for quick executions.
3. **exec family** (e.g., `execl`, `execv`, `execle`, `execve`, etc.):  
Replaces the current process image with a new program. These calls allow a process to execute a different executable, essentially changing its function.
4. **wait:**  
Suspends the calling process until one of its child processes terminates, allowing the parent to retrieve the child's exit status and prevent zombie processes.
5. **waitpid:**  
Similar to `wait`, but it can wait for a specific child process and can be used in a non-blocking manner to avoid being suspended if no children have terminated.
6. **clone:**  
Creates a new process with shared resources, used primarily for threading. This call gives greater control over the shared resources compared to `fork`.
7. **exit:**  
Terminates the calling process and performs necessary cleanup operations. The process can return an exit status to the operating system.
8. **getpid:**  
Returns the process ID of the calling process. This is often used for identifying the current process in various operations.
9. **getppid:**  
Retrieves the parent process ID of the calling process, which can be useful for managing process hierarchies.

10. **setuid:**

Sets the user ID of the calling process. This changes the process's privileges and can enhance security by limiting access.

11. **setgid:**

Sets the group ID of the calling process, influencing group permissions and access rights for resources.

12. **kill:**

Sends a signal to a specified process, which can be used to terminate it or communicate with it. This is crucial for process control and management.

13. **nice:**

Changes the scheduling priority of a process, which can help manage how much CPU time a process receives based on its priority level.

14. **sched\_yield:**

Causes the calling process to relinquish the CPU voluntarily, allowing other processes to run. This can improve responsiveness in multitasking environments.

15. **getpriority:**

Retrieves the scheduling priority of a specified process. This information can help understand the relative importance of processes.

16. **setpriority:**

Adjusts the scheduling priority of a specified process. This can be used to optimize performance based on process importance.

17. **times:**

Provides user and system CPU time consumed by the calling process and its children. This information is useful for performance monitoring.

18. **getrlimit:**

Retrieves resource limits for the calling process, such as maximum memory usage or CPU time. This can help manage resource allocation effectively.

19. **setrlimit:**

Sets resource limits for the calling process, allowing control over how much system resources can be utilized.

20. **ptrace:**

Enables one process to observe and control the execution of another, commonly used for debugging. This can facilitate process management and behavior monitoring.

21. **prctl:**

Provides a way to control the behavior of a process, including setting its name and controlling its child process handling.

22. **sigaction:**

Allows a process to specify how to handle specific signals. This is crucial for managing interruptions and signals from other processes.

23. **sigprocmask:**

Examines and changes the signal mask of the calling process, affecting how signals are handled or blocked.

24. **sigpending:**

Returns the set of signals pending for the calling process. This helps in managing signal delivery and handling.

25. **syscall:**

Provides a generic interface for making system calls. It can be used to invoke system calls directly with specific parameters.

---

## QUESTION#3

This program creates a binary tree of processes, where each process creates two child processes, forming a tree structure of specified depth. Each process prints its own PID (Process ID) and PPID (Parent Process ID). The program also demonstrates zombie processes by killing a last-level parent process using kill(), after which the orphaned child processes print a message indicating they are "zombie" processes.



```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <signal.h>
```

*// Function to print indentation based on the depth level of the process in the tree*

```
void print_indentation(int depth) {
    for (int i = 0; i < depth - 1; i++) {
        printf("  "); // Each level has 4 spaces of indentation
    }
}
```

*// Function to simulate zombie processes if their parent is killed*

```
void check_for_zombie_status() {
    printf("I am a Zombie process (pid %d, ppid %d)\n", getpid(), getppid());
}
```

*// Recursive function to create the binary tree of processes*

```
void create_binary_tree(int depth, int position, int max_depth) {
```

```
    print_indentation(depth);
    printf("[%d] pid %d, ppid %d\n", position, getpid(), getppid());
```

*// **Base case:** if depth is 1, stop creating further child processes*

```
    if (depth == 1) return;
```

```
    pid_t left_pid, right_pid; // Process IDs for left and right child processes
    int left_status, right_status; // Variables to store exit statuses of children
```

*// Fork to create the left child process*

```
    left_pid = fork();
```

```
    if (left_pid < 0) {
        perror("Failed to fork left child");
        exit(1);
    }
```

```
    } else if (left_pid == 0) {
```

*// **Inside the left child process***

```
        create_binary_tree(depth - 1, 2 * position, max_depth); // Recursive call for left child
        exit(position);
```

```
    } else {
```

```

// Back in the parent process, print the creation message for left child
print_indentation(depth);
printf("[%d] pid %d created left child with pid %d\n", position, getpid(), left_pid);

// Wait for the left child to finish before creating the right child
waitpid(left_pid, &left_status, 0);
print_indentation(depth);

printf("[%d] left child %d of %d exited with status %d\n", position, left_pid, getpid(),
WEXITSTATUS(left_status));
}

// Fork to create the right child process

right_pid = fork();

if (right_pid < 0) {
    perror("Failed to fork right child");
    exit(1);
} else if (right_pid == 0)
    // Inside the right child process
    create_binary_tree(depth - 1, 2 * position + 1, max_depth);
    exit(position);
} else {

    // Back in the parent process, print the creation message for right child
    print_indentation(depth);
    printf("[%d] pid %d created right child with pid %d\n", position, getpid(), right_pid);

    // If this is the last level parent (other than main), kill it to create zombie children
    if (depth == max_depth - 1) {
        kill(getpid(), SIGKILL); // Kill this parent process
    }
    waitpid(right_pid, &right_status, 0);
    print_indentation(depth);

    printf("[%d] right child %d of %d exited with status %d\n", position, right_pid, getpid(),
WEXITSTATUS(right_status));
}
}

```

```

int main(int argc, char *argv[]) {
    // Check if the user provided a depth as a command-line argument
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <depth>\n", argv[0]); // Error message if no depth is given
        return 1;
    }
    int depth = atoi(argv[1]);
    if (depth < 1) {
        fprintf(stderr, "Depth must be at least 1.\n");
        return 1;
    }
    printf("Starting process tree with depth %d:\n", depth);
    create_binary_tree(depth, 1, depth);

    // If any child finds its parent killed, it will print the zombie message
    check_for_zombie_status();
    return 0;
}

```

## OUTPUT

```

mariam@mariam-VirtualBox:~/Desktop$ gcc process.c -o process
mariam@mariam-VirtualBox:~/Desktop$ ./process 3
Starting process tree with depth 3:
    [1] pid 11884, ppid 8042
    [1] pid 11884 created left child with pid 11885
    [2] pid 11885, ppid 11884
    [2] pid 11885 created left child with pid 11886
[4] pid 11886, ppid 11885
    [2] left child 11886 of 11885 exited with status 2
    [2] pid 11885 created right child with pid 11887
[5] pid 11887, ppid 11885
    [1] left child 11885 of 11884 exited with status 0
    [1] pid 11884 created right child with pid 11888
    [3] pid 11888, ppid 11884
    [3] pid 11888 created left child with pid 11889
[6] pid 11889, ppid 11888
    [3] left child 11889 of 11888 exited with status 3
    [3] pid 11888 created right child with pid 11890
[7] pid 11890, ppid 1491
    [1] right child 11888 of 11884 exited with status 0
I am a Zombie process (pid 11884, ppid 8042)
mariam@mariam-VirtualBox:~/Desktop$ █

```

